

Peter Kirgis

105 Taylor Court, Princeton NJ 08540
pk7019@princeton.edu — peterkirgis.github.io

Education

Princeton School of Public and International Affairs

May 2025

Master of Public Affairs (MPA), Certificate in Statistics and Machine Learning

- Cumulative GPA: 3.94

Williams College

June 2020

Bachelor of Arts with Honors in Political Economy and Philosophy

Work Experience

Center for Information Technology, Princeton University

August 2025 – Present

AI Research Engineer

- Working with Arvind Narayanan and Zeynep Tufekci on AI evaluation and alignment projects.
- Leading data visualization and analysis for the Holistic Agent Leaderboard (HAL).
- Designed and implemented a data pipeline to extract and grade 100GB of agent logs using Python.

Skills: AI evaluation, AI agent implementation, technical AI writing

Work Products: [HAL Website](#), [Codebase](#)

United States Census Bureau (Coding it Forward)

June 2024 – August 2024

Data Scientist Fellow

- Built sorting algorithms and k -nearest neighbors regression models in Python to improve existing imputation methods within the Economic Statistical Methods Division.
- Designed Python data visualizations comparing imputation strategies across several loss functions.

Skills: Imputation, Non-parametric machine learning, Geospatial analysis with Python

Work Products: [Improving Imputation Accuracy](#)

Commonwealth of Massachusetts

November 2021 – July 2023

Data Analyst

- Led inter-agency data sharing and program evaluation efforts on behalf of the Chief Data Officer.
- Designed and maintained PostgreSQL ETL pipelines and automated data workflows with Python, AWS Lambda, and Tableau.
- Implemented the Commonwealth's first differential privacy algorithm on longitudinal datasets related to early childhood education and workforce development.
- Presented yearly work summary to the Boston Area Research Initiative (Fall 2022).

Skills: SQL, Python, AWS, ETL, Differential privacy, Tableau, PowerBI, Figma, Jira

Work Products: [Employee Diversity Dashboard](#), [Juvenile Justice Data Story](#), [Early Child Integrated Data System](#)

Research Portfolio

Differences in the Moral Foundations of Large Language Models

Princeton, 2025

- Administered a synthetic experiment to sixteen frontier large language models to elicit diverse moral judgments. Used PCA and NLP methods to explore and visualize bias and clustering of LLM responses from each model provider relative to a human baseline.
- Invited talk at the 2025 PICSciE/CSML Joint Colloquium.

[Paper](#) | [Code](#) | [Slides](#)

Desegregation Paradox? A model of LIHTC's effects on economic segregation

Princeton, 2024

- Used novel data sources to simulate the effect of redistributing LIHTC residents across US metropolitan areas and observed that LIHTC's presence is associated with a decrease in economic segregation.

[Paper](#)